

Лабораторная работа №4
по дисциплине
«Методы машинного обучения»
на тему
«Подготовка обучающей и тестовой выборки,
кросс-валидация и подбор гиперпараметров на
примере метода ближайших соседей.»

Выполнил:
студент группы ИУ5-24М
Зубаиров В. А.

```
[54]: import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from sklearn.datasets import load_iris, load_boston
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.model_selection import KFold, RepeatedKFold, LeaveOneOut, LeavePOut,
↳ ShuffleSplit, StratifiedKFold
from sklearn.model_selection import learning_curve, validation_curve
from sklearn.metrics import mean_absolute_error, mean_squared_error,
↳ mean_squared_log_error, median_absolute_error, r2_score
from sklearn.metrics import roc_curve, roc_auc_score
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
```

```
[3]: data = pd.read_csv("heart.csv")
```

```
[4]: data.info(5)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column   Non-Null Count  Dtype
---  ---
0   age      303 non-null    int64
1   sex      303 non-null    int64
2   cp       303 non-null    int64
3   trestbps 303 non-null    int64
4   chol     303 non-null    int64
5   fbs      303 non-null    int64
6   restecg  303 non-null    int64
7   thalach  303 non-null    int64
8   exang    303 non-null    int64
9   oldpeak  303 non-null    float64
10  slope    303 non-null    int64
11  ca       303 non-null    int64
12  thal     303 non-null    int64
13  target   303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
[5]: data.describe()
```

```
[5]:      age      sex      cp  trestbps      chol      fbs \
count 303.000000 303.000000 303.000000 303.000000 303.000000 303.000000 303.000000
mean  54.366337  0.683168  0.966997 131.623762 246.264026 0.148515
std   9.082101  0.466011  1.032052 17.538143 51.830751 0.356198
min   29.000000  0.000000  0.000000 94.000000 126.000000 0.000000
25%   47.500000  0.000000  0.000000 120.000000 211.000000 0.000000
50%   55.000000  1.000000  1.000000 130.000000 240.000000 0.000000
75%   61.000000  1.000000  2.000000 140.000000 274.500000 0.000000
max   77.000000  1.000000  3.000000 200.000000 564.000000 1.000000

      restecg  thalach  exang  oldpeak  slope  ca \
count 303.000000 303.000000 303.000000 303.000000 303.000000 303.000000 303.000000
mean  0.528053 149.646865 0.326733 1.039604 1.399340 0.729373
std   0.525860 22.905161 0.469794 1.161075 0.616226 1.022606
min   0.000000 71.000000 0.000000 0.000000 0.000000 0.000000 0.000000
25%   0.000000 133.500000 0.000000 0.000000 1.000000 0.000000 0.000000
50%   1.000000 153.000000 0.000000 0.800000 1.000000 0.000000 0.000000
75%   1.000000 166.000000 1.000000 1.600000 2.000000 1.000000 1.000000
max   2.000000 202.000000 1.000000 6.200000 2.000000 4.000000

      thal  target
count 303.000000 303.000000
mean  2.313531  0.544554
std   0.612277  0.498835
min   0.000000  0.000000
25%   2.000000  0.000000
50%   2.000000  1.000000
75%   3.000000  1.000000
max   3.000000  1.000000
```

```
[6]: data.corr()
```

```
[6]:      age      sex      cp  trestbps      chol      fbs \
age      1.000000 -0.098447 -0.068653 0.279351 0.213678 0.121308
sex     -0.098447 1.000000 -0.049353 -0.056769 -0.197912 0.045032
cp      -0.068653 -0.049353 1.000000 0.047608 -0.076904 0.094444
trestbps 0.279351 -0.056769 0.047608 1.000000 0.123174 0.177531
chol     0.213678 -0.197912 -0.076904 0.123174 1.000000 0.013294
fbs      0.121308 0.045032 0.094444 0.177531 0.013294 1.000000
restecg -0.116211 -0.058196 0.044421 -0.114103 -0.151040 -0.084189
thalach -0.398522 -0.044020 0.295762 -0.046698 -0.009940 -0.008567
exang    0.096801 0.141664 -0.394280 0.067616 0.067023 0.025665
oldpeak  0.210013 0.096093 -0.149230 0.193216 0.053952 0.005747
slope   -0.168814 -0.030711 0.119717 -0.121475 -0.004038 -0.059894
ca       0.276326 0.118261 -0.181053 0.101389 0.070511 0.137979
thal     0.068001 0.210041 -0.161736 0.062210 0.098803 -0.032019
target  -0.225439 -0.280937 0.433798 -0.144931 -0.085239 -0.028046
```

```

restecg thalach exang oldpeak slope ca \
age -0.116211 -0.398522 0.096801 0.210013 -0.168814 0.276326
sex -0.058196 -0.044020 0.141664 0.096093 -0.030711 0.118261
cp 0.044421 0.295762 -0.394280 -0.149230 0.119717 -0.181053
trestbps -0.114103 -0.046698 0.067616 0.193216 -0.121475 0.101389
chol -0.151040 -0.009940 0.067023 0.053952 -0.004038 0.070511
fbs -0.084189 -0.008567 0.025665 0.005747 -0.059894 0.137979
restecg 1.000000 0.044123 -0.070733 -0.058770 0.093045 -0.072042
thalach 0.044123 1.000000 -0.378812 -0.344187 0.386784 -0.213177
exang -0.070733 -0.378812 1.000000 0.288223 -0.257748 0.115739
oldpeak -0.058770 -0.344187 0.288223 1.000000 -0.577537 0.222682
slope 0.093045 0.386784 -0.257748 -0.577537 1.000000 -0.080155
ca -0.072042 -0.213177 0.115739 0.222682 -0.080155 1.000000
thal -0.011981 -0.096439 0.206754 0.210244 -0.104764 0.151832
target 0.137230 0.421741 -0.436757 -0.430696 0.345877 -0.391724

```

```

thal target
age 0.068001 -0.225439
sex 0.210041 -0.280937
cp -0.161736 0.433798
trestbps 0.062210 -0.144931
chol 0.098803 -0.085239
fbs -0.032019 -0.028046
restecg -0.011981 0.137230
thalach -0.096439 0.421741
exang 0.206754 -0.436757
oldpeak 0.210244 -0.430696
slope -0.104764 0.345877
ca 0.151832 -0.391724
thal 1.000000 -0.344029
target -0.344029 1.000000

```

```
[9]: np.unique(data.target)
```

```
[9]: array([0, 1])
```

```
[15]: target = data.iloc[:, -1]
```

```
[17]: data_data = data.iloc[:, 0:-1]
```

```
[19]: target.shape
```

```
[19]: (303,)
```

```
[20]: data_data.shape
```

```
[20]: (303, 13)
```

```
[21]: heart_X_train, heart_X_test, heart_y_train, heart_y_test = train_test_split(
    data_data, target, test_size=0.5, random_state=1)
```

```
[22]: heart_X_train.shape, heart_y_train.shape
```

[22]: ((151, 13), (151,))

```
[23]: heart_X_test.shape, heart_y_test.shape
```

[23]: ((152, 13), (152,))

```
[24]: np.unique(heart_y_train), np.unique(heart_y_test)
```

[24]: (array([0, 1]), array([0, 1]))

```
[25]: cl1_1 = KNeighborsClassifier(n_neighbors=2)
      cl1_1.fit(heart_X_train, heart_y_train)
      target1_1 = cl1_1.predict(heart_X_test)
      len(target1_1), target1_1
```

[25]: (152,
array([0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0,
 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1,
 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0,
 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,
 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0,
 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1]))

```
[26]: cl1_2 = KNeighborsClassifier(n_neighbors=10)
      cl1_2.fit(heart_X_train, heart_y_train)
      target1_2 = cl1_2.predict(heart_X_test)
      len(target1_2), target1_2
```

[26]: (152,
array([1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1,
 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0,
 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0,
 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0,
 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1]))

```
[27]: accuracy_score(heart_y_test, target1_1)
```

[27]: 0.6052631578947368

```
[28]: accuracy_score(heart_y_test, target1_2)
```

[28]: 0.625

```
[29]: cl1_3 = KNeighborsClassifier(n_neighbors=30)
      cl1_3.fit(heart_X_train, heart_y_train)
      target1_3 = cl1_3.predict(heart_X_test)
      len(target1_3), target1_3
```

```
[29]: (152,
array([[1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
        1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0,
        1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1,
        0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0,
        1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
        1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1,
        0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1,
        0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1]))
```

```
[30]: accuracy_score(heart_y_test, target1_3)
```

```
[30]: 0.618421052631579
```

```
[61]: kf = KFold(n_splits=5)
scores = cross_val_score(KNeighborsClassifier(n_neighbors=10),
                          data_data, target, scoring='f1_weighted',
                          cv=kf)

scores
```

```
[61]: array([0.67391304, 0.67391304, 0.62214834, 0.58823529, 0.63636364])
```

```
[62]: scores = cross_val_score(KNeighborsClassifier(n_neighbors=50),
                              data_data, target,
                              cv=LeaveOneOut())

scores, np.mean(scores)
```

```
[62]: (array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1.,
        0., 1., 1., 1., 1., 1., 0., 1., 1., 1., 0., 0., 1., 1., 0., 1., 0.,
        0., 1., 1., 1., 0., 0., 0., 1., 1., 0., 1., 1., 1., 1., 0., 1., 1.,
        0., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 0., 1., 1., 1., 1.,
        1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0.,
        1., 0., 1., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0., 1., 1., 0.,
        1., 1., 1., 0., 0., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 1.,
        1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 1.,
        0., 1., 1., 0., 1., 1., 1., 1., 0., 0., 0., 1., 1., 1., 0., 0., 1.,
        0., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 0., 0.,
        1., 0., 0., 0., 1., 1., 0., 0., 1., 1., 1., 0., 0., 1., 0., 0.,
        1., 0., 0., 1., 1., 1., 1., 0., 1., 0., 0., 1., 0., 0., 1., 1., 1.,
        1., 0., 1., 0., 0., 0., 0., 1., 0., 1., 1., 1., 1., 1., 1., 0., 0.,
        1., 0., 1., 1., 1., 1., 0., 0., 1., 0., 1., 1., 1., 1., 0., 0., 0.,
        0., 0., 1., 1., 1., 1., 1., 0., 0., 1., 0., 0., 1., 1., 1., 1., 1.,
        1., 1., 1., 0., 0., 0., 0., 1., 0., 1., 0., 1., 0., 1., 1., 1.,
        1., 0., 1., 0., 1., 1., 0., 1., 1., 0., 0., 0., 0., 1., 0., 0., 1.,
        1., 0., 1., 0., 0., 0., 1., 0., 1., 1., 1., 0., 0., 0.]),
0.6633663366336634)
```

```
[ ]:
```


```
[32]: fpr, tpr, thresholds = roc_curve(heart_y_test, target1_2,
                                      pos_label=1)

fpr, tpr, thresholds
```


```
[32]: (array([0.      , 0.42253521, 1.      ]),
      array([0.      , 0.66666667, 1.      ]),
      array([2, 1, 0]))
```

```
[33]: def draw_roc_curve(y_true, y_score, pos_label, average):
      fpr, tpr, thresholds = roc_curve(y_true, y_score,
                                       pos_label=pos_label)
      roc_auc_value = roc_auc_score(y_true, y_score, average=average)
      plt.figure()
      lw = 2
      plt.plot(fpr, tpr, color='darkorange',
               lw=lw, label='ROC curve (area = %0.2f)' % roc_auc_value)
      plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
      plt.xlim([0.0, 1.0])
      plt.ylim([0.0, 1.05])
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('Receiver operating characteristic example')
      plt.legend(loc="lower right")
      plt.show()
```


```
[34]: draw_roc_curve(heart_y_test, target1_1, pos_label=1, average='micro')
```



```
[35]: draw_roc_curve(heart_y_test, target1_2, pos_label=1, average='micro')
```



```
[36]: draw_roc_curve(heart_y_test, target1_3, pos_label=1, average='micro')
```



```
[56]: scores = cross_val_score(KNeighborsClassifier(n_neighbors=5),
                                data_data, target, cv=3)
```

```
[57]: scores
```

```
[57]: array([0.62376238, 0.6039604 , 0.66336634])
```

```
[58]: np.mean(scores)
```

```
[58]: 0.6303630363036303
```

```
[37]: n_range = np.array(range(5,55,5))
      tuned_parameters = [{'n_neighbors': n_range}]
```



```
tuned_parameters
```

```
[37]: [{ 'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50])}]
```

```
[41]: %%time
      clf_gs = GridSearchCV(KNeighborsClassifier(), tuned_parameters, cv=5, scoring='accuracy')
      clf_gs.fit(heart_X_train, heart_y_train)
```

CPU times: user 219 ms, sys: 6.41 ms, total: 225 ms

Wall time: 237 ms

```
[41]: GridSearchCV(cv=5, error_score=nan,
                  estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                  metric='minkowski',
                                                  metric_params=None, n_jobs=None,
                                                  n_neighbors=5, p=2,
                                                  weights='uniform'),
                  iid='deprecated', n_jobs=None,
                  param_grid=[{ 'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40,
45, 50])}],
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='accuracy', verbose=0)
```

```
[42]: clf_gs.best_estimator_
```

```
[42]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                          metric_params=None, n_jobs=None, n_neighbors=35, p=2,
                          weights='uniform')
```

```
[44]: clf_gs.best_score_
```


```
[44]: 0.6823655913978494
```

```
[46]: clf_gs.best_params_
```

```
[46]: { 'n_neighbors': 35}
```

```
[47]: plt.plot(n_range, clf_gs.cv_results_[ 'mean_test_score' ])
```

```
[47]: [<matplotlib.lines.Line2D at 0x12241a5b0>]
```



```
[48]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                             n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):
```


```
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.3,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt
```

```
[50]: plot_learning_curve(KNeighborsClassifier(n_neighbors=14), 'n_neighbors=14',
                             heart_X_train, heart_y_train, cv=20)
```

```
[50]: <module 'matplotlib.pyplot' from
'/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-
packages/matplotlib/pyplot.py'>
```



```
[51]: def plot_validation_curve(estimator, title, X, y,
                                param_name, param_range, cv,
                                scoring="accuracy"):


    train_scores, test_scores = validation_curve(
        estimator, X, y, param_name=param_name, param_range=param_range,
        cv=cv, scoring=scoring, n_jobs=1)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

    plt.title(title)
    plt.xlabel(param_name)
    plt.ylabel(str(scoring))
    plt.ylim(0.0, 1.1)
    lw = 2
    plt.plot(param_range, train_scores_mean, label="Training score",
             color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.4,
                     color="darkorange", lw=lw)
    plt.plot(param_range, test_scores_mean, label="Cross-validation score",
             color="navy", lw=lw)
    plt.fill_between(param_range, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.2,
                     color="navy", lw=lw)
    plt.legend(loc="best")
    return plt
```

```
[52]: plot_validation_curve(KNeighborsClassifier(), 'knn',
                             heart_X_train, heart_y_train,
                             param_name='n_neighbors', param_range=n_range,
```

```
cv=20, scoring="accuracy")
```

```
[52]: <module 'matplotlib.pyplot' from  
'/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-  
packages/matplotlib/pyplot.py'>
```



```
[65]: n_range = np.array(range(5,55,5))  
tuned_parameters = [{'n_neighbors': n_range}]  
tuned_parameters
```

```
[65]: [{'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50])}]
```

```
[67]: %%time  
clf_gs = GridSearchCV(KNeighborsClassifier(), tuned_parameters, cv=LeaveOneOut(),  
    ↪scoring='accuracy')  
clf_gs.fit(data_data, target)
```

CPU times: user 9.65 s, sys: 61.3 ms, total: 9.71 s

Wall time: 9.87 s


```
[67]: GridSearchCV(cv=LeaveOneOut(), error_score=nan,  
    estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,  
    metric='minkowski',  
    metric_params=None, n_jobs=None,  
    n_neighbors=5, p=2,  
    weights='uniform'),  
    iid='deprecated', n_jobs=None,  
    param_grid=[{'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40,  
45, 50])}],  
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
    scoring='accuracy', verbose=0)
```

```
[69]: clf_gs.best_params_
```

```
[69]: {'n_neighbors': 30}
```

```
[70]: plt.plot(n_range, clf_gs.cv_results_['mean_test_score'])
```

```
[70]: [<matplotlib.lines.Line2D at 0x122680820>]
```



```
[ ]:
```

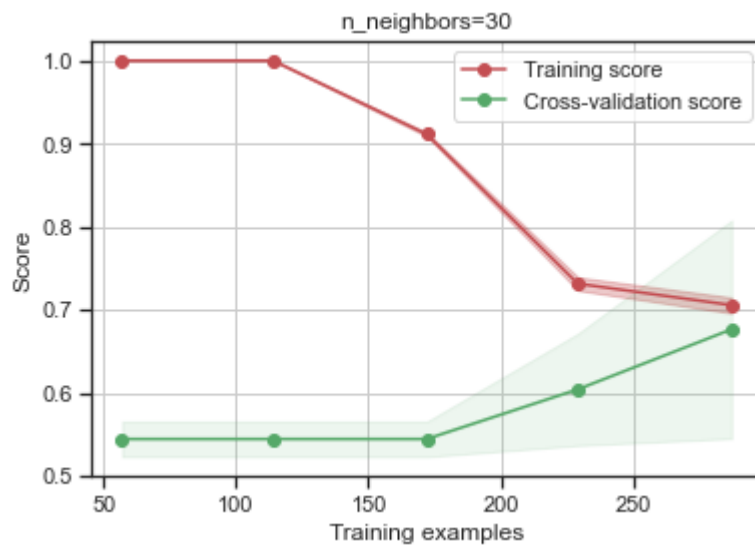
```
[71]: clf_gs.best_estimator_.fit(heart_X_train, heart_y_train)
target2_0 = clf_gs.best_estimator_.predict(heart_X_train)
target2_1 = clf_gs.best_estimator_.predict(heart_X_test)
```

```
[72]: accuracy_score(heart_y_train, target2_0), accuracy_score(heart_y_test, target2_1)
```

```
[72]: (0.7284768211920529, 0.618421052631579)
```

```
[73]: plot_learning_curve(clf_gs.best_estimator_, 'n_neighbors=30',
                        data_data, target, cv=20, train_sizes=np.linspace(.2, 1.0, 5))
```


```
[73]: <module 'matplotlib.pyplot' from
'/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-
packages/matplotlib/pyplot.py'>
```



```
[74]: n_range2 = np.array(range(5,125,5))
```

```
[75]: plot_validation_curve(clf_gs.best_estimator_, 'knn',
    data_data, target,
    param_name='n_neighbors', param_range=n_range2,
    cv=20, scoring="accuracy")
```

```
[75]: <module 'matplotlib.pyplot' from
'/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-
packages/matplotlib/pyplot.py'>
```



```
[ ]:
```