

Лабораторная работа №6
по дисциплине
«Методы машинного обучения»
на тему
«Ансамбли моделей машинного обучения.»

Выполнил:
студент группы ИУ5-24М
Зубаиров В. А.

1. ЛР6. Ансамбли моделей машинного обучения

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, accuracy_score, r2_score
```

```
[2]: data = pd.read_csv("advertising.csv")
```

```
[3]: data.head()
```

```
[3]:   TV Radio Newspaper Sales
1  230.1  37.8    69.2  22.1
2   44.5  39.3    45.1  10.4
3   17.2  45.9    69.3   9.3
4  151.5  41.3    58.5  18.5
5  180.8  10.8    58.4  12.9
```

```
[4]: data_X = data[["TV", "Radio", "Newspaper"]]
```

```
[5]: data_X
```

```
[5]:   TV Radio Newspaper
1  230.1  37.8    69.2
2   44.5  39.3    45.1
3   17.2  45.9    69.3
4  151.5  41.3    58.5
5  180.8  10.8    58.4
..   ...   ...   ...
196  38.2   3.7    13.8
197  94.2   4.9     8.1
198 177.0   9.3     6.4
199 283.6  42.0    66.2
200 232.1   8.6     8.7
```

[200 rows x 3 columns]

```
[6]: data_Y = data[["Sales"]]
```

```
[7]: X_train, X_test, y_train, y_test = train_test_split(
    data_X, data_Y, test_size=0.25, random_state=1)
```

```
[8]: # Качество отдельных моделей
def val_mae(model):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    plt.plot(X_test, y_test, 'g.')
    plt.plot(X_test, y_pred, 'ro')
    plt.show()
    result = mean_absolute_error(y_test, y_pred)
```

```

r2 = r2_score(y_test, y_pred)
print(model)
print('MAE={}'.format(result))
print('R2={}'.format(r2))


```

```

[9]: for model in [
      GradientBoostingRegressor(),
      RandomForestRegressor(n_estimators=50)
    ]:
      val_mae(model)
      print('=====\\n\\n')

```

/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
y = column_or_1d(y, warn=True)



/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:3:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().


This is separate from the ipykernel package so we can avoid doing imports
until

```

GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0, criterion='friedman_mse',
                           init=None, learning_rate=0.1, loss='ls', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=None, subsample=1.0, tol=0.0001,

```

```
validation_fraction=0.1, verbose=0, warm_start=False)
MAE=0.48997309191670874
R2=0.9831579266623767
```



```
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        max_samples=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=50, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
MAE=0.51311999999999992
R2=0.9815468551914713
```

1.1. Модель градиентного бустинга показала лучший результат на тестовой выборке

```
[10]: from sklearn.model_selection import RandomizedSearchCV

n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]

max_features = ['auto', 'sqrt']

max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)

min_samples_split = [2, 5, 10]
```

```

min_samples_leaf = [1, 2, 4]

bootstrap = [True, False]

random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}
random_grid

```

```

[10]: {'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000],
      'max_features': ['auto', 'sqrt'],
      'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None],
      'min_samples_split': [2, 5, 10],
      'min_samples_leaf': [1, 2, 4],
      'bootstrap': [True, False]}

```

```

[11]: rf = RandomForestRegressor()

rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_iter = 100,
                               cv = 3, verbose=2, random_state=42, n_jobs = -1)

rf_random.fit(X_train, y_train)

```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done  9 tasks   | elapsed:  3.3s
[Parallel(n_jobs=-1)]: Done 130 tasks   | elapsed: 20.3s
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 47.8s finished
/usr/local/lib/python3.7/site-packages/sklearn/model_selection/_search.py:739:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
self.best_estimator_.fit(X, y, **fit_params)

```

```

[11]: RandomizedSearchCV(cv=3, error_score=nan,
                        estimator=RandomForestRegressor(bootstrap=True,
                                                           ccp_alpha=0.0,
                                                           criterion='mse',
                                                           max_depth=None,
                                                           max_features='auto',
                                                           max_leaf_nodes=None,
                                                           max_samples=None,
                                                           min_impurity_decrease=0.0,
                                                           min_impurity_split=None,
                                                           min_samples_leaf=1,

```

```

        min_samples_split=2,
        min_weight_fraction_leaf=0.0,
        n_estimators=100,
        n_jobs=None,
oob_score=False,
        param_distributions={'bootstrap': [True, False],
                            'max_depth': [10, 20, 30, 40, 50, 60,
                                           70, 80, 90, 100, 110,
                                           None],
                            'max_features': ['auto', 'sqrt'],
                            'min_samples_leaf': [1, 2, 4],
                            'min_samples_split': [2, 5, 10],
                            'n_estimators': [200, 400, 600, 800,
                                              1000, 1200, 1400, 1600,
                                              1800, 2000]},
        pre_dispatch='2*n_jobs', random_state=42, refit=True,
        return_train_score=False, scoring=None, verbose=2)

```

[12]: rf_random.best_params_

```

{'n_estimators': 800,
 'min_samples_split': 2,
 'min_samples_leaf': 1,
 'max_features': 'auto',
 'max_depth': 100,
 'bootstrap': True}

```

[13]: **def** evaluate(model, test_features, test_labels):

```

    predictions = model.predict(test_features)
    error = mean_absolute_error(y_test, predictions)
    r2 = r2_score(y_test, predictions)
    print('Model Performance')
    print('MAE: {:.4f}'.format(error))
    print('R2 score: {:.4f}'.format(r2))
    print('=====\n\n')

```

```

base_model = RandomForestRegressor(n_estimators = 10, random_state = 42)
base_model.fit(X_train, y_train)
evaluate(base_model, X_test, y_test)

```

Model Performance

MAE: 0.5994

R2 score: 0.9713

=====

/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:11:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using

```
ravel().  
# This is added back by InteractiveShellApp.init_path()
```

```
[14]: best_random = rf_random.best_estimator_  
evaluate(best_random, X_test, y_test)
```

Model Performance

MAE: 0.5178

R2 score: 0.9820

=====

Видно, что подбор гиперпараметров улучшил нашу модель, уменьшив ошибку на 0.08

```
[15]: n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
```

```
max_features = ['auto', 'sqrt']
```

```
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
```

```
max_depth.append(None)
```

```
min_samples_split = [2, 5, 10]
```

```
min_samples_leaf = [1, 2, 4]
```

```
bootstrap = [True, False]
```

```
random_grid_Booster = {'n_estimators': n_estimators,  
                        'max_features': max_features,  
                        'max_depth': max_depth,  
                        'min_samples_split': min_samples_split,  
                        'min_samples_leaf': min_samples_leaf,  
                        }
```

```
[16]: gb = GradientBoostingRegressor()
```

```
gb_random = RandomizedSearchCV(estimator = gb, param_distributions = {  
    ↪ random_grid_Booster, n_iter = 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
```

```
gb_random.fit(X_train, y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.

Fitting 3 folds for each of 100 candidates, totalling 300 fits

[Parallel(n_jobs=-1)]: Done 9 tasks | elapsed: 0.4s

[Parallel(n_jobs=-1)]: Done 221 tasks | elapsed: 7.7s

[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 10.2s finished

/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
[16]: RandomizedSearchCV(cv=3, error_score=nan,
                        estimator=GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0,
                                                            criterion='friedman_mse',
                                                            init=None,
                                                            learning_rate=0.1,
                                                            loss='ls', max_depth=3,
                                                            max_features=None,
                                                            max_leaf_nodes=None,
min_impurity_decrease=0.0,
                                                            min_impurity_split=None,
                                                            min_samples_leaf=1,
                                                            min_samples_split=2,
min_weight_fraction_leaf=0.0,
                                                            n_estimators=100,
                                                            n_...
                        iid='deprecated', n_iter=100, n_jobs=-1,
                        param_distributions={'max_depth': [10, 20, 30, 40, 50, 60,
                                                            70, 80, 90, 100, 110,
                                                            None],
                        'max_features': ['auto', 'sqrt'],
                        'min_samples_leaf': [1, 2, 4],
                        'min_samples_split': [2, 5, 10],
                        'n_estimators': [200, 400, 600, 800,
                                         1000, 1200, 1400, 1600,
                                         1800, 2000]},
                        pre_dispatch='2*n_jobs', random_state=42, refit=True,
                        return_train_score=False, scoring=None, verbose=2)
```

```
[17]: gb_random.best_params_
```

```
[17]: {'n_estimators': 1400,
      'min_samples_split': 10,
      'min_samples_leaf': 2,
      'max_features': 'auto',
      'max_depth': 40}
```

```
[18]: def evaluate(model, test_features, test_labels):
      predictions = model.predict(test_features)
      error = mean_absolute_error(y_test, predictions)
      r2 = r2_score(y_test, predictions)
      print('Model Performance')
      print('MAE: {:.4f}'.format(error))
      print('R2 score: {:.4f}'.format(r2))
      print('=====\\n\\n')

      base_model = GradientBoostingRegressor()
```



```
base_model.fit(X_train, y_train)
evaluate(base_model, X_test, y_test)
```

Model Performance

MAE: 0.4890

R2 score: 0.9832

=====

/usr/local/lib/python3.7/site-packages/sklearn/ensemble/_gb.py:1454:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
y = column_or_1d(y, warn=True)

[19]:

```
best_random = gb_random.best_estimator_  
evaluate(best_random, X_test, y_test)
```

Model Performance

MAE: 0.5078

R2 score: 0.9839

=====

1.2. Подбор параметров в градиентном бустинге не дал прироста качества (оно и так в целом было достаточно высокое)